



Machine Learning Lectures

Explainability And Interpretability

By:
Dr. Andrey Holz

<https://www.linkedin.com/in/andrey-holz>

Explainability and Interpretability



Machine Learning Lectures

Explainability And Interpretability

**By:
Dr. Andrey Holz**


<https://www.linkedin.com/in/andrey-holz>



Agenda

1. Why interpretability matters
2. Interpretability vs. Explainability
3. Global vs. Local Interpretability
4. Intrinsic vs. Post-hoc Interpretability
5. Interpretable Classical Models (Linear, Decision Trees)
6. Feature Importance and Model-Agnostic Methods
7. Advanced Techniques for Explanation
8. Interpreting Specific Classical Models (Ensembles)
9. Limitations and Considerations
10. Key Takeaway

Data Analytics pipeline: CRISP-DM



<https://www.datascience-pm.com/crisp-dm-2/>

<https://web.archive.org/web/20220401041957/https://www.the-modeling-agency.com/crisp-dm.pdf>

Business Understanding	Data Understanding	Data Preparation	Modeling	Evaluation	Deployment
Determine Business Objectives Background Business Objectives Business Success Criteria Assess Situation Inventory of Resources Requirements, Assumptions, and Constraints Risks and Contingencies Terminology Costs and Benefits Determine Data Mining Goals Data Mining Goals Data Mining Success Criteria Produce Project Plan Project Plan Initial Assessment of Tools and Techniques	Collect Initial Data Initial Data Collection Report Describe Data Data Description Report Explore Data Data Exploration Report Verify Data Quality Data Quality Report	Select Data Rationale for Inclusion/Exclusion Clean Data Data Cleaning Report Construct Data Derived Attributes Generated Records Integrate Data Merged Data Format Data Reformatted Data Dataset Dataset Description	Select Modeling Techniques Modeling Technique Modeling Assumptions Generate Test Design Test Design Build Model Parameter Settings Models Model Descriptions Assess Model Model Assessment Revised Parameter Settings	Evaluate Results Assessment of Data Mining Results w.r.t. Business Success Criteria Approved Models Review Process Review of Process Determine Next Steps List of Possible Actions Decision	Plan Deployment Deployment Plan Plan Monitoring and Maintenance Monitoring and Maintenance Plan Produce Final Report Final Report Final Presentation Review Project Experience Documentation

Interpretability vs. Explainability

Interpretability: "level of understanding how the underlying (AI) technology works" [ISO/IEC TR 29119-11:2020(en), 3.1.42]

Directly understand how the model works internally

- **Approach:** Understand how input features lead to outputs
- **Analogy:** Knowing how the car's engine works and what each control does
- **Example:** "If income > \$50,000, approve; otherwise, reject."

Explainability: "level of understanding how the AI-based system ... came up with a given result" [ISO/IEC TR 29119-11:2020(en), 3.1.31]

Explaining model predictions in understandable terms

- **Approach:** Translate complex decisions into accessible narratives
- **Analogy:** Explaining to a passenger why you took a certain turn
- **Example:** Loan approval from a Random Forest: High income (+30%) and good credit score (+50%) help; high debt (-20%) hurts

Why Interpretability and Explainability matter

Trust and Transparency:

Stakeholders need to understand and trust model decisions, especially in critical domains like healthcare, finance, and criminal justice.

Accountability and Ethics:

Ensures models are fair, unbiased, and align with ethical standards in sensitive applications like hiring and credit approvals.

Regulatory Compliance:

Meets legal requirements like GDPR's "Right to Explanation" and industry standards for safety and fairness.

Debugging and Improvement:

Identifies errors, biases, or overfitting to improve performance and reliability.

Actionable Insights:

Enables domain experts to make informed, data-driven decisions by understanding key drivers behind predictions.

When to Use Interpretability vs. Explainability

Audience:

- **Interpretability:** Engineers, regulators who need the technical “inside view” of the model
- **Explainability:** Non-technical stakeholders, customers, decision-makers who need clear, trustworthy reasons behind the model’s output

Use When:

Interpretability:


- Debugging models and ensuring internal logic is solid
- Meeting regulatory requirements for model audits

Explainability:

- Building trust with end-users and clients
- Providing rationale for decisions in contexts like loan approvals or medical diagnoses

Global vs. Local Interpretability

Global Interpretability: Understanding the model's behavior across the entire dataset.




Example:

Knowing that "income" is the most influential feature overall in a decision tree.

Explaining a model to a human decision-maker.

Source: Marco Tulio Ribeiro.

Local Interpretability: Explaining specific predictions for individual instances.



Example:

Explaining why a specific loan was approved based on the applicant's features.

Explaining individual predictions to a human decision-maker.

Source: Marco Tulio Ribeiro.


Intrinsic vs. Post-hoc Interpretability

Intrinsic Interpretability:

Models that are inherently interpretable because of their simplicity and transparency.

- *Examples:* Linear regression, decision trees, rule-based models.
- *Use Case:* Preferred when simplicity and direct understanding are critical, such as in education or low-stakes decisions.

Intrinsic Interpretability: Decision Tree Example




Post-hoc Interpretability:

Methods applied **after the model is trained** to explain its behavior or predictions.

- *Examples:* SHAP, LIME, PDPs.
- *Use Case:* Required for complex models like Random Forests or neural networks to provide explanations.


Post-hoc Interpretability: SHAP Values for an Instance



Feature Importance and Model-Agnostic Methods


- **Permutation Importance:**
 - Measuring drop in performance when a feature is shuffled.
- **Partial Dependence Plots (PDPs):**
 - Visualizing average marginal effects of features.
!features MUST be uncorrelated!
- **Accumulated Local Effects (ALE):**
 - Improved version of PDPs that handle interactions more reliably.

Permutation Importance



How to implement:

- Measure Baseline Performance
- For each feature, shuffle its values while keeping other features unchanged, and recalculate the model's performance
- Compute the drop in accuracy for each feature



Partial Dependence Plots (PDPs):

1-Way PDPs (Left & Middle Panels):

Purpose: Show the isolated effect of a single feature on model predictions.

Temperature (temp): Predictions rise sharply with temperature, peak at moderate levels, and decline at extremes.

Humidity (humidity): Predictions initially improve slightly, but high humidity leads to a sharp decline, showing its negative impact.

2-Way PDP (Right Panel):

Purpose: Visualizes interactions between two features and their joint effect on predictions.

Temperature & Humidity Interaction:

- High temperatures combined with moderate humidity yield the best outcomes (warmer colors).
- High humidity suppresses the positive impact of temperature (cooler colors).

How to calculate PDPs

1. Select the Feature(s):

- Choose the feature (e.g., temperature) or pair of features to analyze.

2. Generate Feature Values:

- Identify the range of values for the feature(s).
- For a pair, create a grid of values to cover both ranges.

3. Replace Feature Values in the Dataset:

- Replace the selected feature's column(s) with each value (or combination of values for 2-way PDPs), keeping all other features constant.

4. Predict for Modified Dataset:

- Use the trained model to predict outcomes for each modified dataset.

5. Average Predictions:



- Calculate the average prediction for each feature value (or combination) to isolate the feature's impact.

6. Plot the Results:

- **1-Way PDP:** Plot feature values on the x-axis and averaged predictions on the y-axis.
- **2-Way PDP:** Use a heatmap or contour plot to visualize interactions.

Accumulated Local Effects (ALE)

- **What ALE Shows:**
 - Captures local feature effects while accounting for interactions.
 - More robust than PDPs for correlated features.
- **Key Steps to Calculate:**
 - Split the feature range into bins.
 - Measure the local effect of changing the feature within each bin.
 - Accumulate the effects across bins.
- **When to Use:**
 - For non-linear models where features interact or are correlated.



Interpretability of Classical Models

Linear Models (Linear/Logistic Regression):


1. **Key Feature:** Coefficients provide a direct understanding of feature influence.
2. **Use Case:** Predicting outcomes where relationships are mostly linear (e.g., housing prices, risk assessment).
3. **Example: Coefficients** in linear regression show how much a unit change in a feature impacts the outcome.

$$\hat{y} = X\hat{\beta}$$

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

Decision Trees:

1. **Key Feature:** Clear, step-by-step decision-making paths.
2. **Use Case:** Classifying or predicting outcomes with interpretable rules (e.g., loan approvals).
3. **Example:** "If income > \$50,000 and credit score > 700, approve the loan."



Interpreting Linear Regression Coefficients

What Linear Regression Coefficients Show

- **Direction of Relationship:**
 - Positive coefficients: Feature increases the predicted outcome.
 - Negative coefficients: Feature decreases the predicted outcome.
- **Magnitude of Impact:**
 - Larger absolute values indicate stronger influence on the prediction.

Key Metrics for Interpretability

1. Coefficients:

Quantify the change in the dependent variable for a one-unit change in the independent variable, holding others constant.

2. P-values:

1. Measure the statistical significance of each coefficient.
2. Low p-values (< 0.05): Indicate statistically significant features

3. Confidence Intervals:

Range within which the true coefficient is likely to fall.

Key Insights from the Regression Summary

Target: MedHouseVal (Median House Value)

Feature	Coefficient	P-value	Interpretation
MedInc	0.359	0	A \$1,000 increase in median income increases house value by 0.359 units.
AveRooms	0.015	0	Each additional average room increases house value by 0.015 units.
AveOccup	-0.0034	0	Higher occupancy slightly reduces house value (-0.0034 per unit increase).
Latitude	-0.498	0	Moving 1 degree north decreases house value by 0.498 units.
Longitude	-0.5117	0	Moving 1 degree west decreases house value by 0.512 units.

Metrics to Note

Statistical Significance ($P > |t|$):

- All features have p-values < 0.05 == statistically significant.

Confidence Intervals:

- **The 95% confidence intervals confirm reliability.**
For example, MedInc is between 0.353 and 0.365.

Condition Number ($1.56e+04$):

- **High condition number suggests potential multicollinearity.**

OLS Regression Results						
=====						
Dep. Variable:	MedHouseVal	R-squared:	0.586			
Model:	OLS	Adj. R-squared:	0.586			
Method:	Least Squares	F-statistic:	5842.			
Date:	Thu, 12 Dec 2024	Prob (F-statistic):	0.00			
Time:	23:05:03	Log-Likelihood:	-23140.			
No. Observations:	20640	AIC:	4.629e+04			
Df Residuals:	20634	BIC:	4.634e+04			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-42.8298	0.647	-66.247	0.000	-44.097	-41.563
MedInc	0.3590	0.003	117.141	0.000	0.353	0.365
AveRooms	0.0150	0.002	6.432	0.000	0.010	0.020
AveOccup	-0.0034	0.000	-6.760	0.000	-0.004	-0.002
Latitude	-0.4980	0.007	-71.939	0.000	-0.512	-0.484
Longitude	-0.5117	0.007	-70.224	0.000	-0.526	-0.497
=====						
Omnibus:	4705.040	Durbin-Watson:	0.729			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	11739.657			
Skew:	1.253	Prob(JB):	0.00			
Kurtosis:	5.715	Cond. No.	1.56e+04			

R-squared: 0.586

- 58.6% of the target variation is explained by the model.
- Indicates a moderately good fit.

F-statistic (5842) & Prob (F-statistic): 0.00

- The model is statistically significant overall.

Interpreting Logistic Regression Coefficients

Logistic regression models the log-odds as a linear function of the features

$$\text{LogOdds} = \ln(\text{Odds}) = X\hat{\beta} \quad \text{Odds} = \frac{1 - P(\text{event})}{P(\text{event})} \quad \beta_i = \Delta \text{LogOdds} / \Delta X_i \quad \text{OddsRatio} = e^{\beta_i}$$

What Linear Regression Coefficients Show

- **Direction of Relationship:**
 - Positive coefficients: Increase the likelihood of the target class.
 - Negative coefficients: Decrease the likelihood of the target class.
- **Magnitude of Impact:**
 - Larger absolute values: Stronger influence on the prediction probability.

- If $\beta_i > 0$: Feature X_i increases the odds.
- If $\beta_i < 0$: Feature X_i decreases the odds.
- If $\beta_i = 0$: Feature X_i has no effect on the odds.

Key Metrics for Interpretability

1. Coefficients:

Represent the change in odds for a one-unit increase in the feature.

2. P-values:

1. Measure the statistical significance of each coefficient.
2. Low p-values (< 0.05): Indicate statistically significant features.

3. Confidence Intervals:

Range where the true odds ratio is likely to fall, ensuring robustness of interpretation.

$$P(\text{event}) = \frac{1}{1 + e^{(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n)}}$$

Key Insights from the Classification Summary

Feature	P	Coefficient	Interpretation	Odds Ratio	Interpretation
const	0	35.003	Baseline log-odds when all features are at 0. Represents a very high baseline probability.		
mean radius	0	-1.161	A 1-unit increase in mean radius decreases the log-odds of the target by 1.161. This feature has a strong negative impact on the target likelihood.	$\exp(-1.161) = 0.31$	A 1-unit increase reduces the odds of the target class by 69%.
mean texture	0	-0.338	A 1-unit increase in mean texture reduces the log-odds of the target by 0.338. This feature also negatively impacts the target probability.	$\exp(-0.338) = 0.71$	A 1-unit increase reduces the odds of the target class by 29%.
mean smoothness	0	-101.533	A 1-unit increase in mean smoothness reduces the log-odds drastically by 101.533, suggesting this feature has a very strong negative influence .	$\exp(-101.533) \approx 0$ (close)	A 1-unit increase almost entirely eliminates the odds of the target class.
mean compactness	0.046	-13.545	A 1-unit increase in mean compactness decreases the log-odds by 13.545. Though significant (p-value = 0.046), its effect is smaller compared to smoothness.	$\exp(-13.545) \approx 0$ (very low)	A 1-unit increase drastically reduces the odds of the target class.

Statistical Significance ($P > |t|$):

- All features have p-values < 0.05 == statistically significant.

Feature Impact:

- mean smoothness**: the largest negative impact (-101.533)
- mean compactness** and **mean radius** also negatively affect the target but to a lesser degree.
- The intercept (**const**) is **positive**, indicating that when all features are at zero, the **baseline log-odds are significantly high**.

Optimization terminated successfully.
Current function value: 0.179271
Iterations 9

Logit Regression Results

Dep. Variable:	target	No. Observations:	398
Model:	Logit	Df Residuals:	393
Method:	MLE	Df Model:	4
Date:	Fri, 13 Dec 2024	Pseudo R-squ.:	0.7289
Time:	01:20:22	Log-Likelihood:	-71.350
converged:	True	LL-Null:	-263.17
Covariance Type:	nonrobust	LLR p-value:	9.483e-82

	coef	std err	z	P> z	[0.025	0.975]
const	35.0030	4.631	7.558	0.000	25.926	44.080
mean radius	-1.1615	0.157	-7.395	0.000	-1.469	-0.854
mean texture	-0.3381	0.062	-5.445	0.000	-0.460	-0.216
mean smoothness	-101.5327	24.138	-4.206	0.000	-148.842	-54.223
mean compactness	-13.5446	6.793	-1.994	0.046	-26.859	-0.230

Pseudo R-squared: 0.7289,

- the model explains about 72.9% of the variability in the data.

- Indicates a moderately good fit.



LLR p-value: 9.483e-82:

- The model is statistically significant overall.

Decision Trees & Feature Importance

Decision Trees rank features based on their importance

- Feature importance is calculated during tree training
- **Example:** Questions for credit scoring



Calculating Feature Importance

Feature importance quantifies the contribution of each feature in reducing impurity across the tree.

1. Impurity Reduction at Each Node:

1. Compute the impurity measure (e.g., Gini impurity or entropy) before and after the split.
2. Calculate the difference, weighted by the proportion of samples reaching that node.

2. Aggregate Across the Tree:

Sum the weighted impurity reductions for each feature across all nodes where the feature is used.

3. Normalization:

Divide by the total impurity reduction to express importance as a proportion, summing up to 1 across all features.

$$FI = \frac{1}{T} \sum_{t=1}^T \sum_{j \in \text{nodes}(t)} \frac{N_j}{N} \Delta I_j$$

- T : Total number of trees (for ensemble methods).
- j : Node index.
- N_j : Number of samples at node jj .
- N : Total number of samples.
- ΔI_j : Change in impurity at node jj .

Interpreting Ensembles with Advanced Techniques

Overview of Ensemble Methods

Ensemble methods combine multiple learning algorithms to achieve better predictive performance than individual models.

Types of Ensemble Methods:

- **Bagging: Reduces variance** by training multiple models on different subsets of data and averaging their predictions.
- **Boosting: Reduces bias** by sequentially training models, each focusing on correcting errors of its predecessors.
- **Stacking:** Combines diverse models by training a meta-model to integrate their predictions.

Explainability in Ensemble Methods

Challenges:

- **Complexity:**
Combining multiple models complicates understanding the overall decision process.
- **Opacity:**
Ensembles often act as "black boxes," making it hard to interpret feature contributions.

Approaches to Enhance Explainability:

- **Model-Agnostic Methods:**
Universal techniques to evaluate feature importance and relationships, applicable to any model.
- **Model-Specific Methods:**
Tailored metrics like Mean Decrease in Impurity (MDI) in Random Forests provide insights into feature significance.

Model-Specific Feature Importance in Ensemble Methods:

Random Forests:

- **Mean Decrease in Impurity (MDI):**
Measures total impurity reduction (e.g., Gini impurity) contributed by each feature across all trees. Higher reductions = greater importance.
- **Mean Decrease in Accuracy (MDA):**
Evaluates feature importance by measuring accuracy loss when feature values are shuffled. Significant accuracy drops = key features.

Gradient Boosting Machines (GBMs):

- **Feature Contribution:**
Assesses importance based on how much each feature reduces the loss function during training. Higher loss reductions = higher importance.

Model-Agnostic Feature Importance in Ensemble Methods:

- **Permutation Feature Importance (PFI):**
 - **How it works:** Shuffle feature values and measure accuracy drop.
 - **Use case:** Identify key features driving ensemble predictions.
- **Partial Dependence Plots (PDPs):**
 - **How it works:** Show how changes in one or two features affect predictions, averaging others.
 - **Use case:** Analyze how ensembles respond to specific features on average.
- **Shapley Values:**
 - **How it works:** Allocate ensemble predictions fairly among features using cooperative game theory.
 - **Use case:** Provide consistent, detailed global and local feature importance.
- **LIME (Local Interpretable Model-Agnostic Explanations)::**
 - **How it works:** Approximate ensemble predictions locally with an interpretable model.
 - **Use case:** Explain individual predictions for ensembles by analyzing feature contributions.

SHAP: SHapley Additive exPlanations



SHAP

<https://shap.readthedocs.io/>

What is SHAP?


- A model-agnostic interpretability method based on cooperative game theory.
- Assigns each feature a contribution value toward a prediction.

Why Use SHAP?

- **Global Insights:** Understand overall feature importance across the dataset.
- **Local Explanations:** Explain individual predictions by attributing contributions to features.

How SHAP Works:

- Considers all possible feature combinations to determine each feature's impact.
- Explains ensemble models like Random Forests and Gradient Boosting Machines effectively.



SHAP Summary Plot (Global Interpretability)


Purpose: Understand overall feature importance across the dataset.

Key Elements:

- **Feature Ranking:** Features are ranked from top (most important) to bottom (least important).
- **SHAP Value (X-Axis):** Indicates the feature's contribution to predictions (positive or negative).
- **Color Coding:** Red indicates high feature values, blue indicates low feature values.
- **Spread:** Shows variability in feature impact across data points.


How to Interpret:

- The wider the spread, the more variable the feature's impact on predictions.
- High SHAP values (far from zero) show stronger influence.
- For **MedInc** (Median Income), higher values (red) push the predictions up, while lower values (blue) pull them down.



SHAP Force Plot (Local Interpretability)

Purpose: Explain individual predictions by visualizing feature contributions.



Key Elements:

- **Base Value:** The model's expected output if no features were considered.
- **Feature Contributions:** Bars push the prediction higher (red) or lower (blue).
- **Output Value:** The final prediction as a sum of the base value and feature contributions.

How to Interpret:

- Positive contributions (red) push the prediction upward; negative (blue) push it downward.
- For example, Features like **AveRooms**, **Longitude**, and **MedInc** push the prediction upward.

LIME (Local Interpretable Model-agnostic Explanations)


Locally approximate complex model behavior with a simple, interpretable model.

- 1.Data Perturbation:** LIME perturbs the input data by creating variations of the original data point.
- 2.Local Model:** It fits an interpretable model (e.g., linear regression) to the perturbed data, approximating the complex model's behavior near the original data point.
- 3.Feature Contribution:** The interpretable model assigns weights to features, highlighting their contribution to the prediction.

Why LIME?

- Focuses on **local interpretability**, explaining how individual predictions are made.
- **Model-agnostic**.

Original paper: <https://arxiv.org/pdf/1602.04938>




How to Read LIME Plots

What the Plot Shows:

- **Feature Contributions:** Displays how each feature affects the prediction for a specific instance.
- **Direction of Impact:** Positive contributions increase the prediction, and negative decrease it.
- **Magnitude of Impact:** The length of each bar indicates the strength of the feature's influence.

How to Interpret LIME Plots

- **Rank Features by Importance:** Features with longer bars have a stronger impact on the prediction.
- **Identify Key Influencers:** Quickly see which features drive the prediction and in which direction.
- **Understand Local Behavior:** Focuses only on the selected instance, making it easier to explain individual predictions.



Live Coding & Demos



Limitations and Considerations

1. Limitations

Model-Specific Challenges

- Trade-off: Interpretability vs. Accuracy
Simpler models may lack predictive power compared to complex black-box models.
- Scalability
Explainers like SHAP and LIME can be computationally expensive for large datasets.

Data and Feature Issues

- Feature Correlation: Correlations can mislead methods like PDPs and SHAP.
- Sparse/Imbalanced Data: Reduces reliability of explainability tools.

Ensemble Complexity

- Harder to interpret combined decisions from models like Random Forests and Gradient Boosting.

Limitations and Considerations

2. Ethical and Practical Considerations

- **Over-reliance on Explanations:**
Simplistic interpretations can cause overconfidence in predictions.
- **Fairness and Bias:**
Tools expose, but do not resolve, biases in data or models.
- **Regulatory Compliance:**
Explanations must meet legal requirements (e.g., GDPR).

3. Model-Agnostic Method Limitations

- **Approximation Accuracy:**
Methods like LIME may oversimplify complex models.
- **Global vs. Local:**
Local interpretability (e.g., SHAP, LIME) doesn't ensure global insights.

Tools and Libraries for Explainability

Tool/Library	Description	Features	Use Case	Link
InterpretML	Open-source Python package for glassbox models and blackbox explainers.	Interpretable models (e.g., EBM). Global and local explanations with interactive visualizations.	Understanding model behavior during training and inference.	InterpretML
SHAP	Based on cooperative game theory to explain individual predictions.	Consistent, locally accurate feature attributions. Supports tree-based and deep learning models.	Effective for both global and local interpretability across diverse models.	SHAP
LIME	Provides local explanations by approximating complex models with interpretable ones.	Model-agnostic, applicable to any classifier or regressor. Explains individual predictions.	Interpreting black-box predictions on a case-by-case basis.	LIME
ELI5	Python library for debugging and interpreting ML models.	Supports scikit-learn, XGBoost, LightGBM, CatBoost. Explains global and individual behavior.	Debugging and understanding ML models with clear visualizations.	ELI5
PiML Toolbox	Integrated Python toolbox for interpretable machine learning.	Supports interpretable models like GAM and GAMI-Net. Model-agnostic explainability tools.	Model development and validation, especially for regulated industries.	PiML Toolbox
Feature Importance (Scikit-learn)	Built-in methods to compute feature importance for tree-based models.	Visualizes feature importance as bar charts. Useful for quick global insights.	Understanding which features have the most impact on model predictions.	Scikit-learn

Key Takeaways

1. Importance of Explainability and Interpretability

- **Why it Matters:** Builds trust, transparency, and ethics in AI for domains like healthcare and finance.
- **Interpretability:** Direct understanding of models (e.g., decision trees, linear models).
- **Explainability:** Post-hoc methods for complex models (e.g., SHAP, LIME).

2. Tools and Techniques

- **Model-Agnostic Methods:** SHAP, LIME, PDPs, Permutation Importance.
- **Model-Specific Techniques:** Feature importance (MDI, MDA), and interpretable models.
- **Libraries:** InterpretML, SHAP, LIME, ELI5, Scikit-learn.

3. Challenges and Considerations

- **Trade-offs:** Balancing interpretability with accuracy.
- **Data Issues:** Correlation, sparsity, imbalance may mislead.
- **Ethics:** Explanations reveal biases but can't fix them; compliance is key.

4. Practical Application

- **Local vs. Global:** SHAP and LIME for instance-level insights; PDPs for global trends.
- **Use Case Alignment:** Match methods to model type and stakeholder needs.

Thank you for
your attention

